

Exact test problems and a work–accuracy protocol for benchmarking nearest correlation matrix solvers

Oyeyinka E. Ibrahim
Department of Industrial and Systems Engineering
Northern Illinois University
ibrahimoyeyinka@gmail.com

Draft of September 13, 2026

Abstract

We specialize the known technique of generating test problems from prescribed optimality conditions to the nearest correlation matrix (NCM) problem, obtaining instances with exactly known primal and dual solutions whose solution rank, near-zero multiplicity and eigenvalue separation are set independently. Using a matched-accuracy protocol measured against the exact solution, we show that conclusions drawn from NCM solver comparisons do not transfer across benchmark regimes. On 270 instances at $n = 100$, the order of SBB-Dual and AGD-SDAJ reverses with the accuracy target and is organized by solution rank. On 270 instances at $n = 500$ neither pattern survives: AGD-SDAJ-BH is cheaper at every target and every rank. Four matrices from the literature, from $n = 94$ to $n = 3250$, are consistent with the same dependence on dimension. The semismooth Newton method with the Borsdorf–Higham line search is the cheapest method under every convention and dimension tested. The dimension-scaled stopping rule $\|\nabla\theta\|_2 \leq 10^{-7}n$ permits roughly 340 times less accuracy at $n = 500$ than at $n = 100$, so a larger problem can appear easier. Finally, a known finite-precision failure of objective-value Armijo tests recurs in a second published method, AGD-SDAJ, where it does not affect the published results, and occurs on two real matrices: at $n = 94$ the naive Newton test costs at least 75 times the work for five fewer correct digits, and at $n = 3250$ unmodified AGD-SDAJ stops at its work cap with error 4×10^{-5} , where the safeguarded variant needs 9 EVDS. It is removed in the Newton method by the published Borsdorf–Higham line search, and in AGD-SDAJ by our adaptation of their safeguard.

Contents

1	Introduction	2
1.1	Related work	4
1.2	Outline	5
2	The KKT-controlled instance family	5
2.1	Dual formulation	5
2.2	Construction and exactness	5
2.3	Prescribed spectral structure	6
2.4	Correlation-like inputs: sufficient conditions for $ G_{ij} \leq 1$	6
2.5	Paired-seed design and screening	8
2.6	Release	8

3	Protocol	8
3.1	Native stopping rules confound cost	8
3.2	Work–accuracy trajectories against the exact solution	8
3.3	Accepted-only accounting	9
3.4	Primitive work, not iterations	9
3.5	Timing protocol	9
3.6	Is the EVD a fair unit of work?	10
3.7	Summary of the protocol	11
4	Two implementation traps	11
4.1	Trap 1: redundant Jacobian factorization	11
4.2	Trap 2: the finite-precision Armijo test	11
4.3	How these findings should be read	12
5	Ranking under different conventions	13
5.1	Solvers	13
5.2	Design	13
5.3	The winner does not change	14
5.4	The ordering below the winner changes with the accuracy target	14
5.5	The ordering varies across the degeneracy grid	15
5.6	The same study at $n = 500$	16
5.7	Four matrices from the literature	19
5.8	Accounting convention	20
5.9	What each solver returns	20
5.10	Native exits and exit reasons	21
6	The dimension-scaled stopping rule	21
7	Reproducibility	23
8	Discussion and limitations	24
A	Proofs for the instance family	26
A.1	Proof of proposition 2.3	26
A.2	Proof of proposition 2.4	26
A.3	Proof of proposition 2.5	26
A.4	Mixed μ	27
B	AGD-SDAJ verification record	27
C	Artifact manifest	29

1 Introduction

Given a symmetric matrix $G \in \mathbb{R}^{n \times n}$, the nearest correlation matrix problem asks for

$$\min_X \frac{1}{2} \|X - G\|_F^2 \quad \text{subject to} \quad X \succeq 0, \text{diag}(X) = e. \quad (1)$$

The problem has a unique solution. Since Higham’s alternating-projections method [15], the literature has produced dual formulations and first-order schemes [22, 7], a quadratically convergent semismooth Newton method [25], its preconditioned and globalized refinement [5], accelerated projection methods [17], and, most recently, accelerated gradient methods with quasi-Newton preconditioning [18]. Most of these papers include numerical comparisons, and those comparisons are how a reader decides which method to use.

This paper is about how such comparisons are made. We ask a narrow question: *when NCM solvers are compared, which methodological choices change the conclusion, and which only change the numbers?* We identify four candidate confounds:

1. *Native stopping rules.* Each method stops when its own criterion is met, so the solvers stop at different accuracies.
2. *Approximate references.* Forward error is measured against the output of another solver instead of against the true solution.
3. *Rejected-iterate accounting.* Line-search trials that are rejected can be credited with reaching a target that the method has not actually reached.
4. *Hidden primitive work.* “Iterations” count different operations in different methods, and some per-iteration work is not recorded at all.

A careful benchmark must control each of these. Without an exact solution, however, the second confound cannot be removed, and without trajectories the first and third cannot be separated.

Contributions. In brief, we give an exact-solution NCM family and a matched-accuracy protocol, and use them to show that conclusions drawn from NCM solver comparisons do not transfer across benchmark regimes: an ordering established at $n = 100$ fails to replicate at $n = 500$, and four matrices from the literature are consistent with the same dependence on dimension. We also show that a known finite-precision line-search defect recurs in a second published method and occurs on real matrices. Each contribution builds on established practice (section 1.1).

1. *An exact-solution NCM family* (section 2). We specialize the standard construction of test problems from prescribed optimality conditions [20, 8] to the NCM problem. For semidefinite programs, the same idea is due to Wei and Wolkowicz [27]. The parts specific to NCM are the projection identity $\Pi_{\mathcal{S}_+^n}(G + \text{Diag } y^*) = X^*$, independent control of rank, strict-complementarity defect and eigenvalue separation, a paired-seed design, sufficient conditions for $|G_{ij}| \leq 1$ together with a fixed-rank saturation result for $\max_{i \neq j} |X_{ij}^*|$, and released instance sets at $n = 100$ and $n = 500$.
2. *A matched-accuracy protocol for NCM* (section 3). We adapt data profiles [24], fixed-target evaluation [14] and work–precision diagrams [13] to spectral solvers. The unit of cost is one EVD, which we check against wall-clock time. Because the reference is exact, two rules become enforceable: rejected trials receive no accuracy credit, and accuracy is matched in forward error. Primitive work is reported separately.
3. *A ranking study with a null headline* (section 5). Five verified solver variants were run on 270 instances at $n = 100$ and 270 at $n = 500$. The cheapest method is the same under all 11 conventions examined. Below it, the ordering changes with the accuracy target and across the grid of cells, organized chiefly by solution rank at $n = 100$; at $n = 500$ that crossover disappears. Only a family that varies one spectral quantity at a time makes such an association visible.

4. *Replication of a known line-search defect* (section 4). The failure of objective-value sufficient-decrease tests near a minimizer is well known [12]. For the NCM dual, Borsdorf analysed it [4, §4.7.1], and this led to the modified step selection of Borsdorf and Higham [5]. The new content is limited to two points. First, the defect recurs in a separately published first-order method [18]. Second, we quantify it on the controlled family: 56/270 stalls for semismooth Newton and 18/270 for AGD-SDAJ, both removed by the Borsdorf–Higham correction. We also document a second, simpler trap (a redundant Jacobian factorization).
5. *An NCM quantification of the stopping-rule confound* (section 6). Dolan, Moré and Munson [10] warned that native stopping tests confound comparisons. For NCM, the widely used rule $\|\nabla\theta\|_2 \leq 10^{-7}n$ leaves the achieved forward error spread over a factor of 383 across solvers.

What we do not claim. We do not claim that any method is best in general. The ranking results are for a synthetic family at $n = 100$ and $n = 500$, with four matrices from the literature (section 5.7) as an external check, all on one machine. SBB-Dual appears here as one of five variants and receives the same exposition as the others. Its theory is developed elsewhere and is not needed for any result in this paper. Its convergence theory is developed in a companion manuscript in preparation [19]; none of the results below depends on it.

1.1 Related work

Benchmarking methodology. Performance profiles [9] compare solvers by the ratio of their cost to the best cost on each problem. Cost is measured at each solver’s own termination, so the solver’s stopping test implicitly sets the accuracy. Dolan, Moré and Munson [10] showed that differing stopping criteria bias such profiles, and proposed rerunning under a common optimality measure. Data profiles [24] count cost in a problem-independent unit and show how rankings change with the required accuracy. COCO [14] uses a ladder of fixed targets against known optima. Work–precision diagrams, standard in the comparison of ODE solvers [13], plot error against a reference solution versus work. Gould and Scott [11] show that performance-profile rankings below the leading solver can change when solvers are added or removed. Beiranvand, Hare and Lucet [3] survey best practice, including known-solution test problems and reproducibility. Our protocol follows these principles; our contribution is to make their exact-reference versions possible for the NCM problem. The second/third-place reversals of section 5 arise with a *fixed* solver set, from the accuracy target and the instance regime, so the mechanism differs from the one studied by Gould and Scott.

Test problems with known solutions. Generating test problems from prescribed optimality conditions is standard for unconstrained optimization [23] and for quadratic programming [20, 8]. Wei and Wolkowicz [27] generate semidefinite programs from a chosen optimal primal–dual pair with prescribed complementarity nullity, and show that nullity correlates with solver difficulty. Our family applies the same idea to the NCM problem. Their generator does not apply directly: it builds a linear-objective SDP and chooses the constraint operator as part of the construction, whereas NCM fixes the constraint operator as $\text{diag}(X) = e$ and leaves only G free. As far as we have found, no NCM study uses instances with a known exact solution. The test sets of [15, 25, 4, 5, 17, 2] are random or data-derived, as are the Anymatrix CORRINV matrices [16] used in [18], and in all of them the reference solutions are computed approximately.

Finite-precision line searches. An objective-value sufficient-decrease test cannot be evaluated accurately near a minimizer, because the difference of function values suffers cancellation. Hager and Zhang [12] introduced approximate Wolfe conditions for this reason. For the NCM dual, Borsdorf [4, §4.7.1] derived the stagnation threshold, $\|\nabla\theta\|$ of the order of the square root of the rounding error in θ (about 10^{-8}), noted that the problem is general to Armijo backtracking, and proposed the alternatives that became the modified step selection of [5]. Section 4 replicates this known phenomenon; it does not rediscover it.

1.2 Outline

Section 2 constructs the instance family. Section 3 states the protocol, section 4 describes the two implementation traps, section 5 reports the ranking study, and section 6 examines the dimension-scaled stopping rule. Section 7 covers reproducibility and section 8 discusses limitations. The appendices contain the proofs for the instance family, the verification record for our reimplementations of AGD-SDAJ, and the artifact manifest.

2 The KKT-controlled instance family

2.1 Dual formulation

Following [22, 25], the Lagrangian dual of (1) is the unconstrained problem

$$\min_{y \in \mathbb{R}^n} \theta(y) := \frac{1}{2} \|\Pi_{\mathcal{S}_+^n}(G + \text{Diag}(y))\|_F^2 - e^\top y, \quad \nabla\theta(y) = \text{diag}(\Pi_{\mathcal{S}_+^n}(G + \text{Diag}(y))) - e. \quad (2)$$

Here $\Pi_{\mathcal{S}_+^n}$ is the projection onto the positive semidefinite cone. The gradient is globally Lipschitz with constant 1. If y^* satisfies $\nabla\theta(y^*) = 0$, then $X^* = \Pi_{\mathcal{S}_+^n}(G + \text{Diag}(y^*))$ is the unique solution of (1). Each evaluation of θ or $\nabla\theta$ requires one symmetric eigendecomposition (EVD) of $C(y) := G + \text{Diag}(y)$.

2.2 Construction and exactness

The idea is to choose the solution first and then build a G for which that solution satisfies the KKT conditions exactly. This is the standard technique of generating instances from prescribed optimality conditions [20, 8], and in particular Wei and Wolkowicz’s generator of semidefinite programs with prescribed complementarity nullity [27], specialized here to the NCM problem. What is specific to NCM is the exact projection identity of proposition 2.2, the independent control of the three spectral quantities in section 2.3, and the entrywise conditions of section 2.4. Their construction does not cover the present problem. It generates the data $(\mathcal{A}, b, C)$ of a *linear*-objective SDP, $\min \text{trace } CX$ subject to $\mathcal{A}(X) = b$, $X \succeq 0$, and is free to choose the constraint operator $\mathcal{A}$ itself (their Algorithm 2.3, Steps 3–4). The NCM problem fixes both the objective, which is quadratic, and the constraint operator, $\mathcal{A} = \text{diag}$ with $b = e$; the only free data is G . The shared element is the idea of prescribing an optimal primal–dual pair and controlling the dimension of its common nullspace, which they call the complementarity nullity and which corresponds to $|\beta|$ here.

Definition 2.1 (KKT-controlled instance). Let $n > r \geq 1$. Draw $U \in \mathbb{R}^{n \times r}$ with independent standard-normal entries and rescale every row to unit Euclidean norm. Set $X^* = UU^\top$, so that $\text{diag}(X^*) = e$, $X^* \succeq 0$ and $\text{rank } X^* = r$ almost surely. Let $Q_0 \in \mathbb{R}^{n \times (n-r)}$ have orthonormal columns spanning $\ker X^*$. For $\mu \in \mathbb{R}_{\geq 0}^{n-r}$ define

$$W = Q_0 \text{Diag}(\mu) Q_0^\top, \quad G = X^* - W + \text{Diag}(\text{diag } W), \quad y^* = -\text{diag}(W). \quad (3)$$

Proposition 2.2 (Exactness). *For every $\mu \geq 0$, $\Pi_{\mathcal{S}_+^n}(G + \text{Diag}(y^*)) = X^*$ and $\nabla\theta(y^*) = 0$. Hence X^* is the unique nearest correlation matrix to G .*

Proof. By construction, $G + \text{Diag}(y^*) = X^* - W$. Both X^* and W are symmetric positive semidefinite, and $X^*W = X^*Q_0 \text{Diag}(\mu)Q_0^\top = 0$ because the columns of Q_0 lie in $\ker X^*$. The ranges of X^* and W are therefore orthogonal. They need not be complementary: if some $\mu_j = 0$, then $\text{range } W$ is a proper subspace of $\ker X^*$. Consequently $X^* - W$ has the spectral decomposition obtained by combining those of X^* and $-W$ on orthogonal subspaces. Its positive part is X^* , and its nonpositive part is $-W$. Thus $\Pi_{\mathcal{S}_+^n}(X^* - W) = X^*$, and $\nabla\theta(y^*) = \text{diag}(X^*) - e = 0$. Since θ is convex and differentiable, y^* minimizes (2). Uniqueness of the primal solution is classical [15, 25]. $\square$

In floating point, the generator recomputes the projection rather than relying on the identity. Over the 270 released instances, the KKT residual is about 1.6×10^{-14} and the exactness error $\|\Pi_{\mathcal{S}_+^n}(C(y^*)) - X^*\|_F$ is about 8.2×10^{-14} .

2.3 Prescribed spectral structure

Because the ranges are orthogonal, the spectrum of $C(y^*) = X^* - W$ is the multiset union

$$\text{spec } C(y^*) = \{\lambda_1(X^*), \dots, \lambda_r(X^*)\} \cup \{-\mu_1, \dots, -\mu_{n-r}\}. \quad (4)$$

In the notation of [25], let α, β, γ index the positive, zero and negative eigenvalues of $C(y^*)$. Then

- $|\alpha| = r$, so the rank of the solution is prescribed;
- $|\beta| = \#\{j : \mu_j = 0\}$ exactly, so the multiplicity of zero eigenvalues at the solution is prescribed;
- an entry $\mu_j = \delta > 0$ produces an eigenvalue $-\delta$, which moves toward the nonsmooth boundary of $\Pi_{\mathcal{S}_+^n}$ as $\delta \rightarrow 0$.

Terminology. Qi and Sun [25] showed that constraint nondegeneracy always holds for the NCM problem, so the family does not, and could not, vary constraint degeneracy. What it varies is *failure of strict complementarity*: zero eigenvalues of $C(y^*)$ ($|\beta| > 0$), and near-zero eigenvalues (small δ) that approach it. Throughout this paper, a *degeneracy cell* means one (r, m, δ) combination in this sense, and the *strict-complementarity defect* means $|\beta|$. The generator uses $\mu = (\mu_{\text{bulk}}, \dots, \mu_{\text{bulk}}, \delta, \dots, \delta)$ with m copies of δ and $\mu_{\text{bulk}} = 0.1$. The separation parameter is $\rho = \delta/\mu_{\text{bulk}}$. The released grid is

$$r \in \{5, 20, 50\}, \quad m \in \{1, 5, 20\}, \quad \delta \in \{10^{-2}, 10^{-4}, 10^{-6}, 10^{-8}, 10^{-10}, 0\},$$

with five paired seeds per rank. This gives $3 \times 3 \times 6 = 54$ cells and 270 instances at $n = 100$. In plots, ρ is shown on a $\log_{10}$ axis and the case $\rho = 0$ is shown separately.

2.4 Correlation-like inputs: sufficient conditions for $|G_{ij}| \leq 1$

Proposition 2.2 places no restriction on the entries of G . For the instances to resemble invalid correlation matrices, which is the setting in which NCM problems arise, the generator additionally requires $\text{diag}(G) = e$ (true by construction), $\lambda_{\min}(G) < 0$, and $|G_{ij}| \leq 1$. The first two are checked directly. The results below explain when the third can be expected to hold. Let P_r be the orthogonal projector onto $\text{range } U$, so that $Q_0Q_0^\top = I - P_r$, and define

$$a := \max_{i \neq j} |X_{ij}^*|, \quad b := \max_{i \neq j} |(P_r)_{ij}|.$$

Table 1: Empirical (a, b) at $n = 100$, 30 seeds per rank. The last column is the mean over seeds of the sufficient bound $(1 - a)/b$ from proposition 2.3.

r	mean a	max a	mean b	max b	mean $(1 - a)/b$
5	0.9905	0.9981	0.0590	0.0708	0.162
20	0.7495	0.8340	0.1483	0.1871	1.690
50	0.5081	0.5719	0.1856	0.2083	2.650

Proposition 2.3 (Deterministic bound, uniform μ). *If $\mu = c e$ with $c > 0$, then $|G_{ij}| \leq a + cb$ for all $i \neq j$. In particular, if $a < 1$, then $0 < c \leq (1 - a)/b$ is sufficient for $|G_{ij}| \leq 1$.*

Proposition 2.4 (Fixed-rank saturation). *Fix $r \geq 2$ and let the rows u_i of U be independent and uniform on S^{r-1} . Then $a \rightarrow 1$ as $n \rightarrow \infty$, along every realization.*

Proposition 2.4 says only that a saturates. It does *not* imply that a fixed $c > 0$ eventually violates proposition 2.3: at fixed r the quantity b also decreases with n , so whether $a + cb$ exceeds 1 depends on the relative rates of $1 - a$ and b , which we do not establish. We therefore make no asymptotic feasibility claim, and the generator relies on entrywise screening.

Proposition 2.5 (High-dimensional regime). *Let the rows of U be independent and uniform on S^{r-1} and let $\delta \in (0, 1)$. Then with probability at least $1 - \delta$,*

$$a \leq \sqrt{\frac{2}{r} \left(2 \log n + \log \frac{2}{\delta} \right)}.$$

In particular $a \rightarrow 0$ whenever $r/\log n \rightarrow \infty$.

The bound follows from the subgaussianity of the uniform distribution on the sphere [26, Thm. 3.4.5] and a union bound; see section A. We state it only for a , the quantity that governs proposition 2.3 in the regime of interest. A companion bound on b would require a model for range U that we do not need here, since the released family lies in the fixed-rank regime of proposition 2.4, where proposition 2.5 does not apply.

Proofs are given in section A. These results give sufficient conditions only. Failure of the bound in proposition 2.3 does not show that some $|G_{ij}|$ exceeds 1, and for that reason the generator checks every entry directly instead of relying on the bound. Proposition 2.3 assumes uniform μ , whereas the released instances use a mixed μ (bulk value 0.1 together with m copies of δ). The extension to mixed μ is left open: the extra off-diagonal term it introduces depends on the choice of orthonormal basis for $\ker X^*$ and is not determined by (n, r, m, δ) alone (section A). The released mixed- μ instances are therefore screened entrywise rather than certified by a bound. None of these results says anything about indefiniteness of G , which is checked separately.

Which regime the released family is in. Table 1 reports empirical values of (a, b) at $n = 100$ over 30 seeds per rank. The family is in the regime of proposition 2.4, not proposition 2.5. The rank is fixed and small relative to n , a is close to 1 at $r = 5$, and the bound on c implied by proposition 2.3 shrinks sharply as the rank decreases.

The table is consistent with the screening outcomes. At $\mu_{\text{bulk}} = 0.1$, the mean sufficient bound at $r = 5$ (0.162) exceeds 0.1, but the worst seeds (max $a = 0.9981$) fall below it. In the paired design (section 2.5), 5 of the 7 candidate seeds were accepted at $r = 5$; the 2 rejections were for $|G_{ij}| > 1$. At $r = 20$ and $r = 50$, 5 of 5 were accepted. With $\mu_{\text{bulk}} = 1$, the sufficient bound fails at every rank tested, and in an earlier design every $r = 5$ candidate was rejected. This is why the generator uses $\mu_{\text{bulk}} = 0.1$.

2.5 Paired-seed design and screening

Each candidate seed is screened once per rank against its entire (m, δ) block. A seed is accepted only if all 3×6 instances built from it pass the checks ($\text{diag}(G) = e$, $|G_{ij}| \leq 1$, $\lambda_{\min}(G) < 0$, and a recomputed KKT residual and exactness error each below 10^{-9}). The accepted seeds are then reused across every (m, δ) cell. Varying the separation therefore changes only μ , with U held fixed, so comparisons along δ are within-seed. Rejections are counted and reported, not silently discarded.

Because of the screening, the $r = 5$ sample is conditioned differently from the $r = 20$ and $r = 50$ samples: at $r = 5$ it excludes the seeds with the largest a . Comparisons across ranks should be read with this in mind.

2.6 Release

We release the 270 paired instances (G , X^* , y^* , SHA-256 hashes of the upper triangle of G , and per-instance screening diagnostics), the generator, and the seed-screening log. Section C lists the files. We also release a paired set at $n = 500$, generated on the same grid. All 270 instances passed screening; seed acceptance was 5 of 7 candidates at $r = 5$ (both rejections for $|G_{ij}| > 1$) and 5 of 5 at $r = 20$ and $r = 50$, with $\max_{i \neq j} |G_{ij}| = 0.99987$ among the accepted $r = 5$ instances. The ranking study of section 5.6 uses all five paired seeds per rank.

3 Protocol

This section sets out the measurement protocol. Each rule addresses a confound that we observed on this family, and each is stated together with the evidence that motivates it.

3.1 Native stopping rules confound cost

Suppose every solver is run to its own exit criterion and the resulting costs are compared. The comparison then credits a method for stopping early as much as for converging quickly. This happens even when all methods share one nominal rule. Under the dimension-scaled rule $\|\nabla\theta\|_2 \leq 10^{-7}n$ used by [18], which at $n = 100$ is 10^{-5} , the six solver variants of section 5 stop at median forward errors between 5.04×10^{-8} and 1.93×10^{-5} , a factor of 383 (section 6, fig. 4). A cost comparison taken at these exits compares different accuracies.

3.2 Work–accuracy trajectories against the exact solution

We take one EVD as the unit of cost. After *every* EVD, whether it produces an iterate, a line-search trial or a Jacobian factorization, we record the pair

$$(\text{cumulative EVDs}, \|X_k - X^*\|_F),$$

together with $\|\nabla\theta\|_2$ and a flag indicating whether the evaluated point was accepted. Here X^* is the exact solution from proposition 2.2, not the output of another solver. All runs use a common, tight tolerance of 10^{-11} and a budget of 4000 EVDs. These trajectories are work–precision diagrams in the sense of [13], and reading costs off them at a ladder of targets follows data profiles [24] and fixed-target evaluation [14]. The difference is that the reference is the exact solution, not the best value found by any solver. For any target ε , the cost can then be read off the trajectory after the run. We define the *cost to ε* as the cumulative EVD count at the first accepted iterate that satisfies

Table 2: Median primitive work at matched tolerance 10^{-11} , 270 instances, $n = 100$. An “outer iteration” denotes a different amount of work in each row.

solver	EVDs	CG iterations	line-search trials	accepted outer its
Newton-SIN-BH	5	24	0	4
Newton-SIN	11	30	0	5
AGD-SDAJ-BH	70	0	14	9
AGD-SDAJ	82	0	25	10
SBB-Dual	46	0	0	46
Dykstra-APM	102	0	0	102

$\|X_k - X^*\|_F \leq \varepsilon$ and continues to satisfy it for every later accepted iterate. A solver that never reaches the target, or reaches it and later loses it, is recorded as not reaching it at that ε .

Trajectories also expose *plateaus*. The cost of a quadratically convergent method takes very few distinct values across many decades of ε : in table 6, Newton-SIN-BH costs 3, 4 or 5 EVDs at every target from 10^{-2} to 10^{-10} . A single-tolerance comparison therefore depends on where the tolerance falls relative to these plateaus.

3.3 Accepted-only accounting

A rejected line-search trial always costs one EVD. Under *accepted-only* accounting it can never be credited with reaching a target, because the method did not move to that point. Under *all-trial* accounting it can. The difference matters. An early version of our own analysis used all-trial accounting and produced crossings between solver curves that disappeared once rejected trials were denied credit. We use accepted-only accounting throughout and report the other convention as a sensitivity check (section 5.8).

3.4 Primitive work, not iterations

Table 2 gives the median primitive work at matched tolerance (10^{-11}) over the 270 instances.

Four accepted Newton iterations use 5 EVDs and 24 Krylov iterations. Forty-six SBB-Dual iterations use 46 EVDs and nothing else. Nine AGD-SDAJ-BH outer iterations use 70 EVDs, because each outer iteration contains $q = 2$ quasi-Newton inner steps, each with its own line search. A table of “iterations to convergence” across these rows would compare four different quantities. We therefore report EVDs, Krylov iterations and line-search trials separately.

3.5 Timing protocol

All timings use Julia with `BLAS.set_num_threads(1)`. A warm-up solve is run and discarded, `GC.gc()` is called before the measured solve, and a single `@elapsed` surrounds the whole solve. A separate instrumented pass (TimerOutputs) gives the breakdown of time. In that breakdown, percentages are taken against the measured elapsed time, so any uninstrumented work shows up as “other” time instead of inflating the EVD share. Timings were measured on all 270 instances at both the native ($10^{-7}n$) and the matched (10^{-11}) tolerance. The trajectory runs are a separate, untimed pass. All timings were measured on one machine: Intel Core 7 150U (10 physical cores, 12 logical), 32 GB RAM, Windows 11 (build 26200), Julia 1.12.6 with OpenBLAS (ILP64) through libblastrampoline. The timing runs pin `BLAS.set_num_threads(1)`; the untimed trajectory pass

Table 3: Microseconds per EVD at matched tolerance (10^{-11}), $n = 100$, single-threaded BLAS.

solver	$\mu\text{s}/\text{EVD}$	relative to cheapest
SBB-Dual	4584	1.00
AGD-SDAJ-BH	4689	1.02
AGD-SDAJ	4730	1.03
Newton-SIN	5262	1.15
Dykstra-APM	5440	1.19
Newton-SIN-BH	6262	1.37




Figure 1: Median wall-clock time per EVD. The spread is a factor of 1.37 and works against Newton, whose EVDs are accompanied by Krylov solves that the EVD count does not record.

at $n = 500$ was run with four BLAS threads, which changes wall-clock only and not the EVD counts it records.

3.6 Is the evd a fair unit of work?

The EVD count is a good proxy for cost only if one EVD costs about the same in every method. Table 3 and fig. 1 give the median wall-clock time per EVD at matched tolerance.

The spread is a factor of 1.37, and the bias works *against* Newton. Each Newton outer iteration includes a Krylov solve that no EVD count records, so counting EVDs understates Newton’s cost. Newton-SIN-BH is cheaper than SBB-Dual by a factor of 9.2 in EVDs (5 against 46) but by a factor of 6.3 in wall-clock time (0.032s against 0.202s). On this family, then, EVD counts are biased in a known direction but preserve rank, and the bias is small compared with the gaps being measured. We report both.

The cost metric does not change the ranking. At matched tolerance, ranking the six variants by median EVDs and by median wall-clock time gives the same order:

$$\text{Newton-SIN-BH} < \text{Newton-SIN} < \text{SBB-Dual} < \text{AGD-SDAJ-BH} < \text{AGD-SDAJ} < \text{Dykstra-APM}.$$

This is a negative control for section 5. The reversals reported there come from the accuracy target and the instance regime. They are not artifacts of using the EVD as the unit of work. The same comparison at $n = 500$ gives a spread of 1.29 (with Dykstra-APM, not Newton, dearest per EVD) and again an identical ranking under the two cost metrics, so this negative control holds at both dimensions.

3.7 Summary of the protocol

1. Measure forward error against an exact X^* .
2. Run every solver to a common tight tolerance and record a trajectory point after every EVD.
3. Read costs off the trajectories at a common target; give rejected trials no accuracy credit.
4. Report EVDS, Krylov iterations, line-search trials and wall-clock time, never iterations alone.
5. Report the feasibility of the returned iterate ($\lambda_{\min}(X)$ and $\|\text{diag}(X) - e\|_\infty$), and state which iterate is returned.
6. Report reach rates beside every median that is computed over a subset of instances.

4 Two implementation traps

This section describes two implementation defects. Neither prevents convergence to the correct solution at the tolerances usually reported, so neither is visible from final errors. Both change measured cost substantially. We present them as evidence for the protocol of section 3. They are not new discoveries.

4.1 Trap 1: redundant Jacobian factorization

A semismooth Newton step at y_k needs the spectral decomposition of $C(y_k)$ to apply the generalized Jacobian [25]. That decomposition was already computed when $\theta(y_k)$ and $\nabla\theta(y_k)$ were evaluated, at the accepted trial of the previous line search. An implementation that recomputes it pays one extra EVD per outer iteration and gains nothing. On the instances where both variants converge, recomputing the decomposition raises Newton’s cost from 7 to 13 EVDS, an increase of about 86% (a factor of 1.86), with an identical final error. Every Newton variant in this paper caches the decomposition. The same class of defect turned up in our first reimplement of AGD-SDAJ (section B, defect 1). It was found in the same way, by counting primitives one at a time against a published count.

4.2 Trap 2: the finite-precision Armijo test

An objective-value sufficient-decrease test,

$$\theta(y_k + \alpha d_k) \leq \theta(y_k) + c \alpha \nabla\theta(y_k)^\top d_k, \quad (5)$$

fails in floating point once the predicted decrease $c \alpha |\nabla\theta(y_k)^\top d_k|$ falls below the rounding error in $\theta(y_k)$. Near the solution $|\nabla\theta^\top d| \approx \|\nabla\theta\|^2$, so the test becomes unresolvable at a gradient norm of roughly $\sqrt{u|\theta|/c}$, where u is the unit roundoff. From that point, no step can be accepted, even a step that improves both the gradient and the forward error. None of this is new. The cancellation problem is well known in general optimization [12]. For the NCM dual, Borsdorf [4, §4.7.1] derived the threshold (about 10^{-8} in $\|\nabla\theta\|$) and proposed alternatives, which became the modified step selection of Borsdorf and Higham [5, Sec. 3.3, Alg. 4.1]. Their Section 3.3 observes that when $|\phi(x)| = 1$, $\|p\|_2 < u^{1/2}$ and $\|\nabla\phi(x)\|_2 < u^{1/2}/2$, the computed values $fl(\phi(x+p))$ and $fl(\phi(x))$ coincide, so that the Armijo condition cannot be verified at all. Their Algorithm 4.1, Step 6, then tests whether the two objective values are *nearly equal*, using $|a-b| < \gamma u(1+|a|+|b|)$ with $\gamma = 100$, and if so accepts the full inexact Newton step provided $\|\nabla f(y_k + \alpha_k d_k)\|_2 / \|\nabla f(y_k)\|_2 \leq 1 - \mu$, falling

Table 4: AGD-SDAJ on `degen-r20-m1-d1e-10-p2`, as published and with the Borsdorf–Higham adaptation.

variant	outer	EVDs	exit	final $\ \nabla\theta\ $	$\ X - X^*\ _F$
AGD-SDAJ (as published)	115	4001	EVD budget	3.60×10^{-9}	6.09×10^{-9}
AGD-SDAJ-BH	10	67	converged	9.40×10^{-12}	1.82×10^{-11}

back to a unit steepest-descent step otherwise. Our implementation uses their parameter values, $\mu = 0.9$, $\rho = 0.5$ and $\sigma = 10^{-4}$.

Semismooth Newton. With the naive test (5), Newton stalls on 56 of the 270 KKT instances, using 55,402 backtracks. With Borsdorf–Higham globalization (parameters $\mu = 0.9$, $\sigma = 10^{-4}$, halving, and $\gamma = 100$ in the near-equality test), it converges on all 270 with 8 backtracks. The count of 56 was reproduced exactly in a separate timing run on a separate code path: at matched tolerance, the naive-Armijo control exits at `max_iter` on 56/270. A predictor based on the ULP margin between the two sides of (5) separates stalling from non-stalling iterations with AUC 0.934. In all 2,181 affected iterations, the full step that was *rejected* would have reduced both the gradient norm and the exact forward error.

AGD-SDAJ. The same failure occurs in AGD-SDAJ [18], a separately published method with a different algorithmic structure. Line 5 of Algorithm 1 in that paper is the objective-value test (5). Run to the common tolerance 10^{-11} , AGD-SDAJ as published exhausts the 4000-EVD budget on 18 of the 270 instances. On instance `degen-r20-m1-d1e-10-p2`, the trajectory contains 3198 rejected monotone line-search trials, with θ and $\|\nabla\theta\|$ fixed at 6.0935×10^{-9} and 3.5986×10^{-9} . The predicted decrease is about 1.3×10^{-21} , well below the floating-point resolution of about 2×10^{-16} . We adapted the Borsdorf–Higham near-equality test to Algorithm 1’s step: when the objective difference cannot be resolved, a trial is accepted if it reduces $\|\nabla\theta\|$. With this change, all 18 stalls disappear (table 4). Across the family, the number of instances reaching the tolerance rises from 252/270 to 270/270, and the worst-case cost falls from 4035 EVDs to 131.

4.3 How these findings should be read

Three points limit what these results show.

1. *The phenomenon is known.* Borsdorf analysed it for the NCM dual [4, §4.7.1], and the fix was published in [5]. What we add is limited to (i) its recurrence in a separately published method that uses a naive Armijo test, where it is invisible under that method’s native stopping rule, and (ii) its quantification on the controlled family (56/270 and 18/270 stalls). For semismooth Newton the remedy is the published Borsdorf–Higham modified line search; for AGD-SDAJ it is our adaptation of their near-equality safeguard to the Algorithm 1 step of that method.
2. *It does not affect Huynh and Hwang’s reported results.* Their stopping rule, $\|\nabla\theta\|_2 \leq 10^{-7}n$, equals 10^{-5} on these instances and stops the method far above the precision floor at which the test fails. The stall appears only when the method is pushed to a tight common tolerance, which is what a matched-accuracy comparison requires. On their own instances, our reimplementations reproduce their published counts exactly for P8 and P9, and matches P7 in outer iterations and residual with one differing backtracking decision (section 7, section B).

3. *The adaptation is ours.* AGD-SDAJ-BH is our modification and is not specified by the original authors. We therefore report it alongside the unmodified method everywhere.

The general point is that whether a method appears to converge at a given tolerance can depend on a line-search detail that is unrelated to the algorithm’s design. A benchmark that stops each method under its own rule never exposes this.

5 Ranking under different conventions

5.1 Solvers

We compare five variants of four methods. All are implemented in Julia with a common evaluation kernel, and every evaluation of θ or $\nabla\theta$ costs exactly one EVD. Unless stated otherwise, each variant starts from $y_0 = 0$ (Dykstra-APM starts from $X_0 = G$). The algorithms below are written at the same level of detail and use the notation $(\theta_k, g_k, X_k) \leftarrow \text{EVAL}(y_k)$ for one EVD of $C(y_k)$, with $X_k = \Pi_{\mathcal{S}_+^n}(C(y_k))$ and $g_k = \text{diag}(X_k) - e$. The stopping criterion in all algorithms is $\|g_k\|_2 \leq \tau$ (Dykstra-APM applies it to the diagonal of its PSD iterate). Here $\tau = 10^{-11}$ unless stated otherwise.

Newton-SIN-BH (algorithm 1) is the semismooth Newton method of Qi and Sun [25]. The Newton system is solved inexactly by conjugate gradients, using the block form of the generalized Jacobian restricted to the positive-eigenvalue columns (with $\Omega_{\beta\beta} = 0$) and a 10^{-12} diagonal shift. The spectral decomposition from the accepted trial is cached (section 4.1), and globalization uses the Borsdorf–Higham modified line search [5, Alg. 4.1, Step 6]. We adopt their globalization only: their Algorithm 4.1 also solves the Newton system with preconditioned MINRES rather than CG, which we do not adopt, so **Newton-SIN-BH** is not a reimplement of their algorithm. *Newton-SIN* is the same method with the naive Armijo test; it appears only as a control in sections 3 and 4.

AGD-SDAJ (algorithm 2) is Algorithm 4 of Huynh and Hwang [18], transcribed from the full text and checked against the published counts (section B). It uses $q = 2$ inner quasi-Newton steps with a diagonal secant preconditioner and the nonmonotone residual line search of their equation (6), followed by one accelerated gradient step with a monotone Armijo line search. *AGD-SDAJ-BH* is our adaptation of that Armijo step, described in section 4.2.

SBB-Dual (algorithm 3) is a spectral (Barzilai–Borwein) gradient method on (2). It uses the BB1 step, clamped to $[\eta, 2 - \eta]$ with $\eta = 0.01$, has no line search, performs one EVD per update, and finishes with a diagonal-rescaling step $\hat{X} = D^{-1/2} X D^{-1/2}$, $D = \text{Diag}(\text{diag } X)$, as in [5].

Dykstra-APM (algorithm 4) is Higham’s alternating projections with Dykstra’s correction [15]. Each iteration uses one EVD. The method returns the PSD half-iterate.

5.2 Design

All five variants were run on the 270 instances of section 2 at $n = 100$, to a common tolerance of 10^{-11} with a budget of 4000 EVDs, and a trajectory was recorded as described in section 3.2. We evaluated 11 conventions: the native exit; a common forward-error target $\|X - X^*\|_F \leq \varepsilon$ for $\varepsilon \in \{10^{-2}, 10^{-4}, 10^{-6}, 10^{-8}, 10^{-10}\}$; and a common dual residual $\|\nabla\theta\|_2 \leq \tau$ for the same five

Table 5: Distinct orderings of median cost across the 11 conventions.

ordering (cheapest first)	conventions
Newton-SIN-BH < SBB-Dual < AGD-SDAJ-BH < AGD-SDAJ < Dykstra-APM	forward $10^{-6}, 10^{-8}, 10^{-10}$; residual $10^{-6}, 10^{-8}, 10^{-10}$; native
Newton-SIN-BH < AGD-SDAJ-BH < AGD-SDAJ < SBB-Dual < Dykstra-APM	forward $10^{-2}, 10^{-4}$; residual $10^{-2}, 10^{-4}$

Table 6: Median EVDS to reach and hold $\|X - X^*\|_F \leq \varepsilon$ (accepted-only). Entries marked * are medians over the subset of instances that reached the target; reach counts are in brackets. All unmarked entries are over 270/270.

ε	Newton-SIN-BH	AGD-SDAJ-BH	AGD-SDAJ	SBB-Dual	Dykstra-APM
10^{-2}	3	6	6	8	15
10^{-4}	4	14	14	16	33.5
10^{-6}	4	28	28	24.5	53.5
10^{-8}	5	45	50* [264]	34	74.5
10^{-10}	5	63.5	72* [253]	43.5	95.5

values. We also compared accepted-only with all-trial accounting. Section 5.6 repeats the study at $n = 500$.

5.3 The winner does not change

Table 5 lists the orderings produced by the 11 conventions. There are two distinct orderings, and Newton-SIN-BH is cheapest in both. *This is a null result for winner reversal.* On this family, the conventions change cost ratios and the order of the lower places. They never change which method is cheapest.

5.4 The ordering below the winner changes with the accuracy target

Tables 6 and 7 and fig. 2 show median cost to a common target. At loose targets ($10^{-2}, 10^{-4}$), both AGD variants are cheaper than SBB-Dual. From 10^{-6} onward, SBB-Dual is cheaper. The same reversal occurs at the same point under the forward-error and the dual-residual criteria, so it does not depend on which error measure is used. A paper that reports at a loose tolerance would place AGD-SDAJ above SBB-Dual, and a paper that reports at a tight tolerance would place them the other way. Both are defensible one-number summaries of the same runs.

Reach and hold. A target counts as reached at the first accepted iterate from which the error stays at or below the target for the rest of the run. The accelerated methods are not monotone in forward error, so this differs from the first iterate at which the error dips below the target; scoring the first dip alone would understate their cost. Under this rule Newton-SIN-BH, AGD-SDAJ-BH, SBB-Dual and Dykstra-APM reach every target on all 270 instances. Only unmodified AGD-SDAJ misses some: 264/270 and 253/270 at forward targets 10^{-8} and 10^{-10} (264 and 254 for the residual), which is consistent with its 18 budget-exhausting stalls (section 4.2). Its entries marked * are medians over the instances it reached, and are therefore optimistic. The second/third-place reversal is the same whether or not those entries are included.




Figure 2: Median cumulative EVDS to reach and hold $\|X_k - X^*\|_F \leq \varepsilon$ (accepted-only accounting) over 270 instances, $n = 100$, on a half-decade grid. SBB-Dual and AGD-SDAJ-BH cross between $\varepsilon \approx 3 \times 10^{-5}$ and 10^{-5} . † AGD-SDAJ-BH and every other solver reach and hold each target shown on all 270 instances. Unmodified AGD-SDAJ does so on all 270 down to 3×10^{-7} , and on 253/270 at 10^{-10} ; below 3×10^{-7} its medians are taken over the instances it reached. The curves stop at 10^{-10} . Cluster-bootstrap intervals for the decade targets are given in table 8; they are narrower than the plotted marks at most targets.

Uncertainty. Table 8 gives cluster-bootstrap 95% intervals for the forward-error medians ($B = 2000$, resampling paired seeds within rank, so that all 18 (m, δ) cells of a drawn seed move together). The intervals are narrow and do not overlap between SBB-Dual and the AGD variants at any target, so the change of order between 10^{-4} and 10^{-6} is not a sampling artifact.

The paired per-instance picture at $n = 100$ is, however, much closer than the aggregate medians suggest. Against AGD-SDAJ-BH at $\varepsilon = 10^{-2}$ the median paired difference is only +2 EVDS (95% CI [1, 2]), SBB-Dual is cheaper on 40.4% of the 270 instances, and 10 instances tie. The $n = 100$ ordering between these two is therefore a modest average effect with substantial per-instance overlap, in contrast to $n = 500$ (section 5.6), where SBB-Dual is cheaper on 0 of 270 instances at the same target.

5.5 The ordering varies across the degeneracy grid

Aggregate medians can hide dependence on the regime. At a fixed target (forward error 10^{-8}), the 54 (r, m, δ) cells produce 7 distinct orderings (table 9).

Newton-SIN-BH is cheapest in every cell. Below it, at $n = 100$, the order is governed almost entirely by the rank r of the solution. AGD-SDAJ-BH is cheaper than SBB-Dual in all 18 cells with $r = 5$; SBB-Dual is cheaper in 16 of the 18 cells with $r = 20$ and in all 18 with $r = 50$. Within a rank, the near-zero multiplicity m and the separation δ rarely change this pair’s order. Both AGD variants fall below Dykstra-APM to last place in 14 cells: all 12 cells with $r = 50$ and $m \in \{5, 20\}$, and the two cells with $r = 20$, $m = 20$ and $\delta \in \{10^{-8}, 10^{-10}\}$. This block is determined by rank and multiplicity; the separation plays no visible role in it. The instance family was built to vary these quantities one at a time. On a collection of real matrices with uncontrolled spectra, the ordering

Table 7: Median EVDs to reach and hold $\|\nabla\theta\|_2 \leq \tau$ (accepted-only), with reach counts as in table 6.

τ	Newton-SIN-BH	AGD-SDAJ-BH	AGD-SDAJ	SBB-Dual	Dykstra-APM
10^{-2}	2	5	5	7	12
10^{-4}	3	13	13	14	30
10^{-6}	4	27	27	23	50
10^{-8}	5	42	47* [264]	32.5	71
10^{-10}	5	60	69* [254]	41.5	92

Table 8: Cluster-bootstrap 95% intervals for the medians of table 6 ($n = 100$, $B = 2000$, resampling the five paired seeds within each rank). Bracketed counts are instances reaching and holding the target.

ε	Newton-SIN-BH	AGD-SDAJ-BH	AGD-SDAJ	SBB-Dual	Dykstra-APM
10^{-2}	3 [3, 3]	6 [6, 6.5]	6 [6, 6.5]	8 [8, 8]	15 [15, 15]
10^{-4}	4 [3, 4]	14 [13, 15]	14 [13, 15]	16 [15.5, 16]	33.5 [33.5, 33.5]
10^{-6}	4 [4, 4]	28 [27, 29]	28 [27, 29]	24.5 [24.5, 24.5]	53.5 [53.5, 53.5]
10^{-8}	5 [5, 5]	45 [42, 48.5]	50 [46, 52] [264]	34 [34, 34]	74.5 [74.5, 74.5]
10^{-10}	5 [5, 5]	63.5 [62, 66]	72 [70, 75] [253]	43.5 [43, 43.5]	95.5 [95.5, 95.5]

could not be associated with rank in this way.

The attribution is nevertheless specific to $n = 100$. Section 5.6 repeats the study at $n = 500$, where SBB-Dual is no longer cheaper at any rank. What survives across both dimensions is the direction of the effect, not the crossover: SBB-Dual gains on the AGD variants as rank increases.

That rankings depend on the required accuracy is expected in principle [24, 14]. What the controlled family adds is that the variation is organized by a single spectral quantity which the construction sets directly.

Each cell contains only five instances, so a per-cell median is a noisy statistic, and some of the rarer orderings may reflect sampling variation rather than a real change of regime. Resampling the five paired seeds of each cell ($B = 2000$) shows that the count of distinct orderings overstates how much genuinely varies. Most cells reproduce their observed ordering in nearly every resample, but the adjacent pair AGD-SDAJ-BH against AGD-SDAJ is usually *unresolved*: the 95% interval of its median paired difference contains zero in most cells, which is unsurprising because the two differ only in a line-search safeguard that rarely activates at $n = 100$. Reading unresolved adjacent pairs as ties collapses much of the variety in table 9. What survives is the placement of SBB-Dual relative to the AGD pair, which is resolved in the great majority of cells and is the comparison discussed above. Figure 3 shows the same comparison cell by cell as a ratio of median costs, which separates the three controlled quantities: rank sets the level, the near-zero multiplicity modulates it within a rank, and the separation has almost no effect.

5.6 The same study at $n = 500$

The instance family was regenerated at $n = 500$ with the same grid. All 270 instances screened successfully; the acceptance rate at $r = 5$ fell to 5 of 7 candidate seeds, the first rank to show rejections, which is what proposition 2.4 predicts. The ranking study uses all five paired seeds of each rank, 270 instances, with caps of 1200 iterations for SBB-Dual and 800 for Dykstra-APM. It

Table 9: Orderings by cell at forward target 10^{-8} (54 cells, five paired instances each), with the number of cells of each rank producing each ordering.

ordering (cheapest first)	cells	$r=5$	$r=20$	$r=50$
Newton-SIN-BH < SBB-Dual < AGD-SDAJ-BH < AGD-SDAJ < Dykstra-APM	19	0	13	6
Newton-SIN-BH < AGD-SDAJ-BH < AGD-SDAJ < SBB-Dual < Dykstra-APM	18	17	1	0
Newton-SIN-BH < SBB-Dual < Dykstra-APM < AGD-SDAJ-BH < AGD-SDAJ	13	0	2	11
Newton-SIN-BH < AGD-SDAJ-BH < SBB-Dual < AGD-SDAJ < Dykstra-APM	1	0	1	0
Newton-SIN-BH < SBB-Dual < AGD-SDAJ-BH < Dykstra-APM < AGD-SDAJ	1	0	1	0
Newton-SIN-BH < AGD-SDAJ < AGD-SDAJ-BH < SBB-Dual < Dykstra-APM	1	1	0	0
Newton-SIN-BH < SBB-Dual < Dykstra-APM < AGD-SDAJ < AGD-SDAJ-BH	1	0	0	1

Table 10: $n = 500$, 270 instances. Median EVDs to reach and hold $\|X - X^*\|_F \leq \epsilon$, with cluster-bootstrap 95% intervals ($B = 2000$, resampling paired seeds within rank) and reach counts where these are below 270.

ϵ	Newton-SIN-BH	AGD-SDAJ-BH	AGD-SDAJ	SBB-Dual	Dykstra-APM [‡]
10^{-2}	3 [3, 3]	9 [9, 9]	9 [9, 9]	44 [44, 44]	87 [86.5, 88]
10^{-4}	4 [3, 4]	30 [29, 32]	30 [29, 32]	97 [97, 97]	195 [194.5, 195.5]
10^{-6}	4 [4, 4]	53.5 [50.5, 56]	55.5 [52, 58] ^[268]	152 [152, 152]	191 [191, 192] ^[180]
10^{-8}	5 [5, 5]	79 [76, 80]	87 [84, 89] ^[259]	207.5 [207, 208]	258 [258, 260] ^[180]
10^{-10}	5 [5, 5]	101.5 [98.5, 105]	119 [117, 123] ^[242]	260 [259, 262] ^[261]	325.5 [325, 327.5] ^[180]

was run in two parts, seeds 0–2 (about 8.6 h) and seeds 3–4 (about 7.1 h), with identical code and settings.

Two of the $n = 100$ conclusions survive and two do not.

Confirmed. Newton-SIN-BH is again cheapest at every target, at 5 EVDs. The stall of unmodified AGD-SDAJ recurs (31 of 270 runs exhaust the EVD budget, against none for AGD-SDAJ-BH, whose worst case is 240 EVDs against 4039).

Not reproduced. The change of order with the accuracy target seen in table 6 is absent: at $n = 500$ a single ordering holds at all five targets, with AGD-SDAJ-BH ahead of SBB-Dual throughout. The paired per-instance comparison is one-sided. SBB-Dual is cheaper on 0 of 270 instances at 10^{-2} , 3 at 10^{-4} and 15 at 10^{-8} . We report these counts and the cluster-bootstrap intervals of table 10 rather than an instance-level significance test: the design reuses each drawn U across all 18 (m, δ) cells of its rank, so the 270 comparisons form 15 seed clusters, and a test treating them as independent would overstate the evidence by a wide margin. The rank effect of section 5.5 keeps its direction but no longer produces a crossover: the median paired difference at 10^{-8} falls from +669.5 EVDs at $r = 5$ to +129 at $r = 20$ and +13 at $r = 50$, where SBB-Dual is cheaper on 17% of instances. Cell by cell, SBB-Dual is ahead of AGD-SDAJ-BH in 0 of the 54 cells with a resolved bootstrap interval (47 resolved in favour of AGD-SDAJ-BH, 7 unresolved). SBB-Dual is thus behind at every rank tested at $n = 500$, while it led at $r \geq 20$ at $n = 100$.

This is not an artifact of the caps. SBB-Dual converged on all 270 instances (worst case 1075




Figure 3: Ratio of median EVDs to reach and hold $\|X - X^*\|_F \leq 10^{-8}$, SBB-Dual divided by AGD-SDAJ-BH, for each of the 54 cells at $n = 100$ (five paired instances per cell). Values above 1 mean AGD-SDAJ-BH is cheaper; below 1, SBB-Dual is. The ratio is organized primarily by rank: it exceeds 1 in every $r = 5$ cell and is below 1 in every $r = 50$ cell. The near-zero multiplicity m modulates the level within a rank, most visibly at $r = 20$ (from 1.03 at $m = 1$ to 0.40 at $m = 20$), whereas the separation ρ has little effect, as the near-constancy along each row shows.

EVDS against its 1200 cap) and reached every target except 10^{-10} on nine instances. [‡]Dykstra-APM, by contrast, *is* censored: 90 of 270 runs stop at the 800-iteration cap, and from 10^{-6} down its medians cover only 180 instances. Its position at $n = 500$ should be read as censored rather than measured.

The scaling is the cleaner way to state the difference. Between $n = 100$ and $n = 500$ the median cost to 10^{-8} rises by a factor of 6.1 for SBB-Dual (34 to 207.5 EVDS) and by 1.8 for AGD-SDAJ-BH (45 to 79), while Newton-SIN-BH is unchanged at 5. The order below the winner therefore reflects the test set, including its dimension, rather than a fixed property of the methods — which is the paper’s point, in a stronger form than the $n = 100$ data alone supports.

Table 11: Real matrices, target 10^{-9} . EVDS to reach and hold 10^{-8} , total EVDS, measured time at the matched tolerance (smaller two matrices only) and final error. [‡]Stops at its cap and is censored.

matrix	solver	EVDS to 10^{-8}	total EVDS	time (s)	final error
Rocky Mountain $n = 94$	Newton-SIN-BH	4	4	0.009	1.49×10^{-13}
	SBB-Dual	7	8	0.011	1.53×10^{-10}
	Dykstra-APM	11	13	0.023	2.57×10^{-10}
	AGD-SDAJ-BH	12	17	0.023	3.47×10^{-10}
	AGD-SDAJ	12	25	0.036	3.47×10^{-10}
cor1399 $n = 1399$	Newton-SIN-BH	7	7	7.00	1.27×10^{-9}
	AGD-SDAJ-BH	66	91	65.00	6.92×10^{-10}
	AGD-SDAJ	67	136	100.04	2.29×10^{-9}
	SBB-Dual	200	212	152.69	1.89×10^{-9}
	Dykstra-APM [‡]	406	300	288.90	1.49×10^{-6}
cor3120 $n = 3120$	Newton-SIN-BH	6	6	–	2.33×10^{-9}
	AGD-SDAJ-BH	41	42	–	1.32×10^{-9}
	AGD-SDAJ	56	59	–	2.17×10^{-10}
	SBB-Dual	255	270	–	3.38×10^{-9}
	Dykstra-APM [‡]	–	300	–	1.16×10^{-5}
bccd16 $n = 3250$	Newton-SIN-BH	4	5	–	2.71×10^{-11}
	SBB-Dual	5	6	–	1.20×10^{-10}
	Dykstra-APM	6	7	–	1.46×10^{-10}
	AGD-SDAJ-BH	8	9	–	8.88×10^{-10}
	AGD-SDAJ [‡]	–	323	–	3.93×10^{-5}

5.7 Four matrices from the literature

The family of section 2 is synthetic. To check that the regime dependence above is not an artifact of the construction, we ran the same protocol on four invalid correlation matrices from the literature: the Rocky Mountain carbon-storage matrix ($n = 94$), **cor1399** ($n = 1399$), **cor3120** ($n = 3120$) and **bccd16** ($n = 3250$), all from the Anymatrix **CORRINV** collection [16] used by [18].

No exact solution is available here, which is the situation the rest of this paper avoids. Forward error is measured against a reference computed by semismooth Newton driven towards 10^{-13} ; for **cor1399** that reference solve itself stopped at its iteration limit after 1770 EVDS with $\|\nabla\theta\|_2 = 4.59 \times 10^{-13}$. At $n \approx 3000$ an absolute tolerance of 10^{-13} is not attainable in double precision, so for **cor3120** and **bccd16** the reference solve is capped at 30 EVDS and its achieved residual is reported instead: $\|\nabla\theta\|_2 = 1.58 \times 10^{-12}$ and 8.07×10^{-13} , after 2201 s and 895 s. We therefore treat these runs as a check on external validity, not as measurements of the same quality as table 6. The common target is 10^{-9} ; each solver is capped at 500 EVDS on the two smaller matrices and at 300 iterations on the two larger ones. Wall-clock time was measured only on the two smaller matrices.

The matrices disagree with each other, broadly in the direction the synthetic family predicts. At $n = 94$, SBB-Dual is the cheapest of the first-order methods, 7 EVDS to 10^{-8} against 12 for AGD-SDAJ-BH and 11 for Dykstra-APM, which is the pattern of table 9 at $n = 100$. At $n = 1399$, SBB-Dual costs three times AGD-SDAJ-BH (200 against 66), and at $n = 3120$ six times (255 against 41), which is the pattern of section 5.6. **bccd16** is easy for every method: all but one reach 10^{-8} within 8 EVDS, so it does not discriminate between the first-order methods. Newton-SIN-BH is cheapest on all four, as everywhere else in this paper. The change of order is therefore visible on matrices with uncontrolled spectra that we did not construct, and on the two smaller matrices the

Table 12: Accepted-only vs all-trial accounting (forward error, reach and hold). Only rows where the median changed are shown; bracketed counts are instances reaching the target.

ε	solver	accepted-only	all-trial	Δ
10^{-4}	AGD-SDAJ-BH	14.0	15.0	+1.0
10^{-4}	AGD-SDAJ	14.0	15.0	+1.0
10^{-8}	AGD-SDAJ-BH	45.0	46.5	+1.5
10^{-8}	AGD-SDAJ	50.0 [264]	52.0 [264]	+2.0
10^{-10}	AGD-SDAJ-BH	63.5	65.5	+2.0
10^{-10}	AGD-SDAJ	72.0 [253]	74.0 [253]	+2.0

same ordering holds under wall-clock time. With four matrices that differ in more than dimension, this is consistent with the dimension effect of section 5.6, not an independent demonstration of it.

The line-search stall occurs on real data. On the Rocky Mountain matrix the naive-Armijo Newton control exhausts its budget at a forward error of 1.03×10^{-8} , while the same solver with the Borsdorf–Higham line search converges in 4 EVDs to 1.49×10^{-13} . On `cor1399` the naive control converges normally (13 EVDs against 7). On `bccd16` the same defect appears in unmodified AGD-SDAJ. Both AGD variants follow an identical path to a forward error of 3.93×10^{-5} at the fifth EVD. From there AGD-SDAJ-BH, whose safeguard accepts the step, converges in four more EVDs to 8.9×10^{-10} ; unmodified AGD-SDAJ rejects trial after trial in its monotone line search, with the error unchanged in its first eight digits, until it exhausts the budget at 323 EVDs. The defect of section 4.2 is thus not an artifact of the constructed family, and it affects both published methods on real data, but neither is it universal: whether it appears depends on the matrix.

The cost of a reference. At $n \approx 3000$ the reference solve alone took 15 to 37 minutes, and a first attempt at `bccd16` with an uncapped absolute tolerance ran for four hours without finishing. A reference accurate enough to measure forward error can cost more than the solves being measured. That is a practical argument for the synthetic family, whose solution is available in closed form at any size.

5.8 Accounting convention

Under all-trial accounting, every evaluated point counts toward reaching and holding a target, including rejected line-search trials. A rejected trial can then be credited with reaching a target, but one that lies above the target also breaks a hold. On this family the second effect dominates slightly. Median costs change only for the two AGD variants (table 12), by at most 2 EVDs, and all-trial accounting is marginally the more expensive. Newton-SIN-BH, SBB-Dual and Dykstra-APM are unaffected, because they have few or no rejected trials here. The effect is negligible on this family and never changes an ordering. For methods with many rejected trials it need not be, and accepted-only accounting is the rule that cannot credit a point the method itself rejected.

5.9 What each solver returns

Methods also differ in what they return (table 13). All five return matrices with $\lambda_{\min}$ at roundoff level, so none is meaningfully more positive semidefinite than another. The diagonal is a different matter. SBB-Dual returns an exactly unit diagonal because of its rescaling epilogue, whereas every

Table 13: Feasibility of the returned iterate at matched tolerance 10^{-11} , 270 instances. “diag err” is $\|\text{diag}(X) - e\|_\infty$.

solver	median $\lambda_{\min}(X)$	worst $\lambda_{\min}$	median diag err	worst diag err
Newton-SIN-BH	-3.42×10^{-15}	-1.25×10^{-14}	8.38×10^{-15}	5.70×10^{-12}
AGD-SDAJ-BH	-3.40×10^{-15}	-1.36×10^{-14}	2.53×10^{-12}	7.20×10^{-12}
AGD-SDAJ	-3.40×10^{-15}	-1.39×10^{-14}	2.79×10^{-12}	4.01×10^{-8}
SBB-Dual	-3.42×10^{-15}	-1.49×10^{-14}	0	0
Dykstra-APM	-3.43×10^{-15}	-1.36×10^{-14}	1.49×10^{-12}	9.37×10^{-12}

Table 14: Cost at the exit of each solver under the common tolerance 10^{-11} .

solver	median EVDs	min	max	median final $\ X - X^*\ _F$
Newton-SIN-BH	5	4	10	1.25×10^{-13}
AGD-SDAJ-BH	70	27	131	1.10×10^{-11}
AGD-SDAJ	82	34	4035	1.16×10^{-11}
SBB-Dual	46	13	218	1.74×10^{-11}
Dykstra-APM	102	24	443	1.98×10^{-11}

other variant returns a diagonal that is accurate only to its convergence tolerance. The worst-case diagonal error of unmodified AGD-SDAJ is 4.01×10^{-8} , about 2.5 orders of magnitude larger than that of any other variant, which is a direct consequence of the stalls in section 4.2. The rescaling epilogue has a cost of its own: on a nearly feasible input at loose tolerance, it can increase $\|X - X^*\|_F$. Borsdorf and Higham bound that increase by $\text{tol} \|X\|_F / (1 - \text{tol})$ [5, Lem. 3.2], which is small whenever the tolerance is tight but need not be at a loose one. A benchmark that reports only cost hides these differences. A user who needs an exactly unit diagonal and a user who needs minimal distance may not be best served by the same method.

5.10 Native exits and exit reasons

For completeness, table 14 reports cost at each solver’s own exit when all are run to 10^{-11} , and table 15 the exit reasons. Under this convention, the ordering equals the tight-target ordering in table 5, because the common tolerance is tight.

6 The dimension-scaled stopping rule

A common exit criterion for dual NCM methods is $\|\nabla\theta(y_k)\|_2 \leq 10^{-7}n$. It is used, for example, in the experiments of [18]. Because the threshold grows linearly with n , the forward error it guarantees depends on both the dimension and the method. The general point that native stopping tests confound comparisons was made by Dolan, Moré and Munson [10]. This section documents a specific NCM instance of it.

Spread across methods at a fixed n . Table 16 and fig. 4 report the timing run under this rule at $n = 100$, where the tolerance is 10^{-5} .

At this dimension the best and worst median achieved accuracies differ by a factor of 383. That factor is specific to $n = 100$: repeating the same run at $n = 500$ (18 instances, threshold 5×10^{-5})

Table 15: Exit reasons, 270 instances. “at x^{pre} ” denotes the early exit at the preconditioned point of [18].

solver	exit	count
Newton-SIN-BH	converged	270
AGD-SDAJ-BH	tolerance met / at x^{pre}	177 / 93
AGD-SDAJ	tolerance met / at x^{pre} / EVD budget	156 / 96 / 18
SBB-Dual	tolerance met	270
Dykstra-APM	tolerance met	270

Table 16: Solvers exiting when the stopping quantity specific to each method first reaches the common numerical threshold $10^{-7}n$ at $n = 100$ (for the dual methods this quantity is $\|\nabla\theta\|_2$; for Dykstra-APM it is the diagonal error of the PSD iterate), 270 instances, single-threaded timing.

solver	median EVDs	median time (s)	median $\ X - X^*\ _F$
Newton-SIN-BH	4	0.023	5.04×10^{-8}
Newton-SIN	7	0.035	5.04×10^{-8}
AGD-SDAJ-BH	20	0.085	9.58×10^{-6}
AGD-SDAJ	20	0.084	9.58×10^{-6}
SBB-Dual	18	0.088	7.73×10^{-6}
Dykstra-APM	40	0.222	1.93×10^{-5}

gives a spread of only 13.8, with medians from 1.74×10^{-5} for Newton-SIN-BH to 2.41×10^{-4} for Dykstra-APM. The spread across solvers is therefore not itself a stable quantity, and we do not claim it as one. What does grow with n is the inaccuracy the rule permits, which is the subject of the next paragraph. Newton overshoots the tolerance because a quadratically convergent step moves far past the threshold. First-order methods stop close to it, and Dykstra-APM’s stopping quantity (the diagonal error of its PSD iterate) is not the dual gradient norm at all. The cost column is therefore a comparison at different accuracies, and it favours methods that stop early. At this tolerance the native costs also place SBB-Dual (18 EVDs) slightly ahead of both AGD variants (20 EVDs). That comparison sits at the crossover of fig. 2: the accuracies reached, near 10^{-5} , lie where the matched-accuracy order changes, so the native ranking agrees with neither the loose-target nor the tight-target regime and should not be read as evidence for either.

Growth with n . The rule loosens as n grows. In earlier runs, outside the ranking study, a nearly feasible $n = 500$ instance exited after *zero* updates with forward error 7.8×10^{-4} . At $n = 2000$, both dual first-order methods stopped after 3 to 5 iterations with forward errors between 2.7×10^{-3} and 4.1×10^{-3} . Under this rule, a larger problem can appear easier simply because less accuracy is being asked for. The paired family makes the same point with the dimension controlled. Under this rule Newton-SIN-BH exits at a median forward error of 5.04×10^{-8} at $n = 100$ but 1.74×10^{-5} at $n = 500$, about 340 times worse, purely because the threshold is proportional to n . Every other variant degrades similarly. A reader comparing published $n = 100$ and $n = 500$ results under this rule is comparing solutions of very different quality. The real-matrix runs of section 5.7 were made at a matched tolerance only, so they give no native-rule spread on real data.




Figure 4: Achieved accuracy when each solver exits under the same published dimension-scaled rule (tolerance 10^{-5} at $n = 100$). Median forward errors range from 5.04×10^{-8} (Newton variants) to 1.93×10^{-5} (Dykstra-APM), a factor of 383. Costs measured at these exits are therefore costs at different accuracies.

Table 17: AGD-SDAJ ($q = 2$): published counts (outer iterations / EVDs / nonmonotone backtracks / $\|A - X^*\|_F$) against our reimplementations. Published values are from Table 3 and Table A1 of [18].

instance	published	this work	status
P9 <code>bccd16</code> , $n = 3250$	1 / 5 / 0 / 29.06	1 / 5 / 0 / 29.0563	exact
P8 <code>cor3120</code> , $n = 3120$	3 / 18 / 1 / 5.44	3 / 18 / 1 / 5.4437	exact
P7 <code>cor1399</code> , $n = 1399$	4 / 31 / 8 / 21.03	4 / 32 / 7 / 21.0339	one backtrack differs

Recommendation. Timing or iteration counts should never be reported without the achieved forward error (or, when X^* is unknown, a certified bound on it). Where a dimension-scaled rule is used for compatibility with earlier work, it should be accompanied by a run to a common tight tolerance.

7 Reproducibility

Independent implementations of the earlier solvers. SBB-Dual, semismooth Newton and Dykstra-APM were implemented independently in Python and in Julia, using *different* constructions of the reference solution. On all 14 shared instances, the two implementations agree on update counts and forward errors to every printed digit.

AGD-SDAJ: verified against the published counts. AGD-SDAJ exists only in Julia. Because this paper argues that implementation details can reverse conclusions, we did not report any result for our reimplementations on the KKT family until it had reproduced Huynh and Hwang’s own published counts on their own instances. Those instances are three matrices from their test set, loaded under the same convention used elsewhere in this project (table 17).

P9 and P8 agree exactly. P7 agrees on outer iterations and residual norm, and differs by one backtracking decision: 7 rejected trials locally against 8 published. That single decision accounts for the whole difference in EVDs. Excluding backtracks, the local base cost is $32 - 7 = 25$ and the published base cost is $31 - 8 = 23$. The gap of 2 EVDs is exactly the early-termination saving described in the note to their Table 3: with the extra backtrack, their trajectory reaches

the tolerance *at* the preconditioned point and skips the final gradient step, whereas ours reaches it one step later. We attribute the differing decision to floating-point differences. The reference implementation is in MATLAB and ours in Julia. The two call different LAPACK eigensolver drivers, and on a dense 1399×1399 problem, differences in the last bits of the spectrum are enough to flip a single acceptance test whose two sides are nearly equal. We did not search over parameters to make P7 match, because tuning to a target would be exactly the kind of hidden fitting this paper warns against. Section B gives the full verification record, including three transcription defects found and corrected along the way.

Single-machine caveat. All timings come from one machine with single-threaded BLAS: Intel Core 7 150U (10 physical cores, 12 logical), 32 GB RAM, Windows 11 (build 26200), Julia 1.12.6 with OpenBLAS (ILP64) through libblastrampoline. Per-EVD ratios (table 3) may differ on other hardware and BLAS libraries, and with multithreading. Each configuration was timed once, after a discarded warm-up, so we do not report run-to-run variance.

Artifact. The generator, the instances with hashes, all solver code, the trajectory and timing CSV files, and the analysis scripts that produce every table in this paper are listed in section C. They are available at <https://github.com/Realosunboy6/ncm-exact-benchmark>. An archived snapshot carrying a permanent identifier will be deposited on acceptance.

8 Discussion and limitations

What changed the conclusions, and what did not. On this family, none of the conventions we tested changed which method was cheapest. The accuracy target, the instance regime and the problem dimension did change the order below it: the change of lower-place ordering with the accuracy target that we found at $n = 100$ does not reproduce at $n = 500$ (section 5.6), and the rank effect that organizes it at $n = 100$ no longer changes the order there. The stopping rule changed the achieved accuracy by a factor of 383. The cost unit (EVDs against wall-clock time) and the accounting rule (accepted-only against all-trial) changed cost ratios but no ordering. Two line-search details changed whether methods converged at all. We think the correct summary is conservative. The protocol matters mainly for the *magnitude* of reported advantages and for the ordering of methods that are close in cost. On the evidence here, it does not matter for identifying a clear winner. We expect this to hold wherever one method is ahead by a large factor; where methods are within a factor of about two, conclusions drawn under a single convention should be treated with caution.

Why an exact family. The regime dependence in section 5.5 is organized mainly by rank, one of the three quantities the construction controls. Finding it requires both an exact X^* , so that forward error is not contaminated by an approximate reference, and independent control of the spectral structure at the solution. Real matrices supply neither. The same machinery then showed the limit of that finding: at $n = 500$ (section 5.6) the rank effect keeps its direction but no longer changes the order, so what looked like a rule about the methods at one dimension turns out to depend on the benchmark regime as well.

Limitations.

- *Problem size.* Results are reported at $n = 100$ and $n = 500$ (270 instances each). The $r \geq 20$ crossover between SBB-Dual and AGD-SDAJ found at $n = 100$ does *not* persist at $n = 500$,

where AGD-SDAJ-BH is cheaper at every rank. No dimension above 500 was tested on the family. Dykstra-APM is censored at $n = 500$ by its iteration cap.

- *Synthetic family.* The KKT family is synthetic. Its G matrices are correlation-like but are not estimated from data. Section 5.7 adds four matrices from the literature, from $n = 94$ to $n = 3250$, measured against a computed rather than an exact reference; for the two largest that reference is capped and reaches $\|\nabla\theta\|_2 \approx 10^{-12}$ rather than 10^{-13} , and no wall-clock times were taken. Four matrices cannot separate dimension from the other ways in which they differ.
- *Statistical uncertainty.* Aggregate medians at both dimensions carry cluster-bootstrap intervals (tables 8 and 10). Each dimension rests on five paired seeds per rank, that is 15 seed clusters, so inferential statements remain modest. Cell-level orderings rest on five instances each, and within a cell the adjacent pair AGD-SDAJ-BH against AGD-SDAJ is usually unresolved, so the count of distinct orderings should not be read as a count of real regime changes. We report counts and intervals rather than significance tests, because the reuse of each drawn U across all cells of its rank makes instance-level comparisons strongly dependent.
- *Single machine.* All timings come from one machine with single-threaded BLAS.
- *Our adaptation.* The Borsdorf–Higham adaptation to AGD-SDAJ is ours, not the original authors’. It is reported beside the unmodified method throughout.
- *Feasibility propositions.* Proposition 2.3 assumes uniform μ , and the extension to mixed μ is open (section A). Proposition 2.5 is stated for a only; no companion bound on b is given, since the released family lies in the fixed-rank regime where proposition 2.5 does not apply.
- *Scope of methods.* We benchmark five variants of four methods. Other methods (e.g. Anderson-accelerated alternating projections [17]) are not included. Neither are the recent semismooth Newton method for general projection equations [2] or the factor-structured variant [6].
- *Attribution of the line search.* We use the Borsdorf–Higham globalization but not the rest of their algorithm, and the adaptation of their near-equality safeguard to AGD-SDAJ is ours (sections 4.2 and 5.1).
- *Unweighted problem.* Only the unweighted Frobenius problem is considered.

Conclusion. An exact instance family with a prescribed solution structure makes it possible to separate stopping-rule, reference, accounting and primitive-work effects in NCM benchmarks. On 270 instances at $n = 100$, these effects changed cost ratios and the lower places of the ranking, but not the cheapest method. Repeating the study on 270 instances at $n = 500$ sharpened that finding: Newton-SIN-BH remains the leading method, but the change of lower-place ordering with the accuracy target does not reproduce, and the rank effect that organizes it at $n = 100$ no longer changes the order. Rankings below the leading method are therefore not intrinsic properties of the solvers; they depend materially on the benchmark regime, including its dimension. The one clear failure we found is a known finite-precision line-search defect, which appears in two separately published methods, on the constructed family and on real matrices. It is removed in the semismooth Newton method by the published Borsdorf–Higham modified line search, and in AGD-SDAJ by our adaptation of their near-equality safeguard. We recommend that NCM comparisons report forward

error against a trusted solution at a common tolerance, give rejected trials no accuracy credit, and report primitive work and wall-clock time together.

A Proofs for the instance family

Throughout, $U \in \mathbb{R}^{n \times r}$ has rows u_i , $X^* = UU^\top$, P_r is the orthogonal projector onto $\text{range } U = \text{range } X^*$, and $Q_0 Q_0^\top = I - P_r$ (because X^* is symmetric, $\ker X^*$ is the orthogonal complement of $\text{range } X^*$).

A.1 Proof of proposition 2.3

With $\mu = ce$, we have $W = cQ_0 Q_0^\top = c(I - P_r)$. For $i \neq j$, the diagonal correction $\text{Diag}(\text{diag } W)$ in (3) does not affect the entry (i, j) , and $I_{ij} = 0$, so

$$G_{ij} = X_{ij}^* - W_{ij} = X_{ij}^* + c(P_r)_{ij}.$$

By the triangle inequality, $|G_{ij}| \leq |X_{ij}^*| + c|(P_r)_{ij}| \leq a + cb$. If $a < 1$ and $c \leq (1 - a)/b$, then $a + cb \leq 1$. $\square$

Remark A.1. The bound is sufficient but not necessary. Cancellation between X_{ij}^* and $c(P_r)_{ij}$, and the fact that a and b are generally attained at different pairs (i, j) , both make it conservative. For this reason the generator checks $|G_{ij}| \leq 1$ entry by entry.

A.2 Proof of proposition 2.4

Fix $r \geq 2$ and $\epsilon \in (0, 1)$. Because S^{r-1} is compact, it can be covered by finitely many caps $C_1, \dots, C_N$, $N = N(\epsilon)$, each of angular radius $\epsilon/2$. Any two points in the same cap make an angle of at most ϵ , so their inner product is at least $\cos \epsilon \geq 1 - \epsilon^2/2$. Once $n > N$, the pigeonhole principle forces two of $u_1, \dots, u_n$ into a common cap, so $a \geq 1 - \epsilon^2/2$ deterministically. Since ϵ is arbitrary and $a \leq 1$ always, $a \rightarrow 1$ along every realization. (A probabilistic argument is needed only for a *rate*; the limit itself follows from pigeonhole.) $\square$

Remark A.2 (What this does not give). It is tempting to conclude that a fixed $c > 0$ must eventually violate proposition 2.3. That does not follow. At fixed r the projector bound b also tends to zero, so $a + cb \leq 1$ can persist: with $a_n = 1 - 2/n$ and $b_n = 1/n$ one has $a_n \rightarrow 1$ while $a_n + cb_n < 1$ for every n at $c = 1$. A claim in either direction needs the relative rates of $1 - a$ and b , which we do not establish here.

A.3 Proof of proposition 2.5

Fix $i \neq j$ and condition on u_i . Since u_j is uniform on S^{r-1} and independent of u_i , the uniform distribution on the sphere is subgaussian: by [26, Thm. 3.4.5], for every unit vector v and every $t \geq 0$,

$$\Pr\{u_j^\top v \geq t\} \leq 2 \exp(-t^2 r/2).$$

Applying this with $v = u_i$, and again to $-u_j$, gives $\Pr\{|u_i^\top u_j| \geq t\} \leq 4 \exp(-t^2 r/2)$; the bound is unchanged after averaging over u_i . Taking a union bound over the fewer than $n^2/2$ unordered pairs,

$$\Pr\{a \geq t\} \leq 2n^2 \exp(-t^2 r/2).$$

Setting the right-hand side equal to δ and solving for t gives $t = \sqrt{(2/r)(2 \log n + \log(2/\delta))}$, which is the stated bound. If $r/\log n \rightarrow \infty$ then $t \rightarrow 0$ for any fixed δ . $\square$

Remark A.3 (On b). We do not state a companion bound for $b = \max_{i \neq j} |(P_r)_{ij}|$. Writing $P_r = VV^\top$ with V having orthonormal columns, $(P_r)_{ij}$ is an inner product of two rows of V , whose joint distribution is not that of two independent uniform vectors; a bound would need a model for range U (for instance a Haar-random subspace) together with a justification that the row-normalized Gaussian U induces it. Since the released family lies in the fixed-rank regime of proposition 2.4, where proposition 2.5 does not apply, and since the generator screens $|G_{ij}| \leq 1$ entrywise in any case, we leave this open.

Remark A.4 (Scope). These propositions concern $|G_{ij}| \leq 1$ only. Indefiniteness of G ($\lambda_{\min}(G) < 0$) is not implied and is checked separately. The released family ($r \in \{5, 20, 50\}$, $n = 100$) is in the fixed-rank regime of proposition 2.4, not the high-dimensional regime of proposition 2.5.

A.4 Mixed μ

For the released $\mu = (\mu_{\text{bulk}} e_{k-m}, \delta e_m)$, write $Q_0 = [Q_b \ Q_\delta]$. Then $W = \mu_{\text{bulk}} Q_b Q_b^\top + \delta Q_\delta Q_\delta^\top = \mu_{\text{bulk}}(I - P_r) - (\mu_{\text{bulk}} - \delta) Q_\delta Q_\delta^\top$. The off-diagonal entries therefore pick up an extra term $(\mu_{\text{bulk}} - \delta)(Q_\delta Q_\delta^\top)_{ij}$, whose size depends on the choice of basis for $\ker X^*$ used to select Q_δ . A bound analogous to proposition 2.3 follows with b replaced by $b + \max_{i \neq j} |(Q_\delta Q_\delta^\top)_{ij}|$.

Remark A.5 (Open). We do not characterize $\max_{i \neq j} |(Q_\delta Q_\delta^\top)_{ij}|$. It depends on which orthonormal basis of $\ker X^*$ is chosen, and the generator fixes that basis through the eigendecomposition it computes, so the quantity is not determined by (n, r, m, δ) alone. A basis-free bound would need either a canonical choice of Q_δ or a bound uniform over admissible bases. We therefore state proposition 2.3 for uniform μ only, and the generator screens the mixed- μ instances entrywise rather than relying on a bound. This is why the released family is screened rather than certified.

B AGD-SDAJ verification record

Gate. Before running our reimplement of AGD-SDAJ [18] on the KKT family, we required it to reproduce the authors' own published counts on their own instances. An unverified reimplement that is then reported as losing would not be a defensible result.

Source and implementation. The source is Huynh and Hwang (2025), *Numer. Linear Algebra Appl.* 32(6), DOI 10.1002/nla.70045. We transcribed Algorithm 1 (AGD), Algorithm 3 (QN-SDAJ), Algorithm 4 (AGD-SDAJ) and the Section 4.3 parameters from the full text: $c = 10^{-4}$, $\rho = 0.5$, $\lambda_1 = \lambda_2 = 0.01$, $\epsilon_l = 10^{-5}$, $\epsilon_u = 10^8$, $q = 2$, $x_0 = 0$, the stopping rule $\|\nabla\theta\| \leq 10^{-7}n$, and $k_{\max} = 200$. The implementation is `agd_sda_j.jl` (standalone, used for verification) and `agd_sda_j_ncm` in `bench_sbb_dual.jl` (instrumented, used for the study). The instances are the Anymatrix CORRINV files already used elsewhere in this project, exported with the same loading convention.

Equation (6), as transcribed. The nonmonotone line search of Algorithm 3 is a Li–Fukushima-type rule [21] on the residual:

$$\|\nabla\theta(x + \alpha d)\|^2 \leq (1 + \epsilon_k) \|\nabla\theta(x)\|^2 - \lambda_1 \|\alpha \nabla\theta(x)\|^2 - \lambda_2 \|\alpha d\|^2, \quad \epsilon_k = \frac{1}{(k+1)^2}, \quad (6)$$

with $\lambda_1 = \lambda_2 = 0.01$ and $\alpha = \rho^m$ backtracked from 1. The merit function is $\frac{1}{2} \|\nabla\theta\|^2$. The quasi-Newton direction $d = -B^{-1} \nabla\theta$ need not be a descent direction for it, which is why a nonmonotone

rule is needed. The paper also confirms that B_k is reset to the identity once x_k^{pre} is generated, and that Step 4 discards x_k^{pre} if it does not decrease θ .

Before the full text was available, we tried two guesses: a Grippo–Lampariello–Lucidi rule on θ , and a Li–Fukushima rule with unsquared norms and $\epsilon_k = 0$. Both had the right general form but three wrong details: unsquared norms, a missing $(1 + \epsilon_k)$ slack, and λ_1, λ_2 attached to the wrong terms. Because the two guesses bracketed the published counts, they helped locate the remaining gap.

Table 18: Outer iterations / EVDS / nonmonotone backtracks, and $\|A - X^*\|_F$, under the final transcription (eq6) and the two superseded guesses.

instance	published	eq6 (final)	GLL guess	LF guess
P9 bccd16	1/5/0, 29.06	1/5/0, 29.0563 (exact)	1/5/0	1/5/0
P8 cor3120	3/18/1, 5.44	3/18/1, 5.4437 (exact)	3/17/0	3/19/2
P7 cor1399	4/31/8, 21.03	4/32/7, 21.0339	6/40/3	5/48/17

Results. Under every variant, $\|A - X^*\|_F$ matched the published value to the displayed precision on all three instances. This showed early on that the objective, gradient, projection and stopping rule were correct, and that the discrepancy lay in the amount of work per iteration.

Three hazards when reimplementing Algorithm 4. Each of the following is easy to introduce when transcribing the method from its pseudocode. Each leaves the method converging to the correct solution while changing the work it does, so none is visible from final errors, only from work counts checked against the published ones. We record them for anyone reimplementing the method. All three occurred in our own first transcription and were found by exactly that check: the first two by accounting for the five published EVDS on P9 one at a time, the third because a gap in outer iterations on P7 remained once the line search was correct.

1. *Redundant EVD at an already-evaluated point.* The inner QN-SDAJ routine recomputed θ and $\nabla\theta$ at the point its caller had just evaluated. This is the same class of defect as section 4.1, and it was found in the same way. Fix: pass the caller’s values into the inner routine.
2. *Missing early termination at the preconditioned point.* The note to Table 3 of [18] states that the loop also stops whenever x_k^{pre} satisfies the stopping criterion, which saves fewer than two EVDS. Without this check, the outer gradient step is always paid.
3. *Wrong acceleration factor in Algorithm 3.* The factor β_k^{pre} had been implemented with Algorithm 1’s $a_k = \|g\|^2$. That formula is the special case for $d = -\nabla\theta$. In Algorithm 3, $d = -B^{-1}g$, so the general form of their Section 2.3 (the acceleration of Andrei [1]) is needed:

$$\beta = \frac{-g^\top d}{(g_y - g)^\top d} \quad \text{if the denominator is positive, else } \beta = 1.$$

This form reduces exactly to Algorithm 1’s when $d = -g$, so it is consistent with both statements in the paper. *This defect caused the gap in outer iterations:* correcting it reduced P7 from 6 outer iterations to 4.

The cumulative effect on P7 was 46 EVDS $\rightarrow$ 40 (defects 1–2) $\rightarrow$ 43 (exact equation (6)) $\rightarrow$ 32 (correct acceleration factor), against a published 31.

The remaining P7 residual. See section 7. P7 matches in outer iterations (4) and residual norm (21.0339) and differs by one backtracking decision (7 against 8). That decision accounts for the whole difference in EVDs, through the early-termination saving. We attribute it to last-bit differences between the MATLAB and Julia eigensolvers. P8 and P9 reproduce every count exactly on the same code path. No further parameter search was carried out.

Conditions of use (all met in this paper). (i) The code and this record are published with the results. (ii) The P7 residual is disclosed. (iii) Published counts are printed beside local counts wherever the comparison appears (tables 17 and 18).

The Borsdorf–Higham adaptation. AGD-SDAJ-BH changes only the monotone Armijo test of Algorithm 1, line 5 (algorithm 2). When $|\theta_t - \theta| < \gamma u(1 + |\theta_t| + |\theta|)$ with $\gamma = 100$ and $u = 2^{-53}$, a trial that reduces $\|\nabla\theta\|$ is accepted. The nonmonotone rule (6) is unchanged. This adaptation is ours.

C Artifact manifest

All paths are relative to `experiments/bench/` in the project repository. The generator, solver code, instance manifests, result files and analysis scripts are available at <https://github.com/Realosunboy6/ncm-exact-benchmark>. Instance binaries are not stored there: they reach 1 GB at $n = 500$ and are exactly reproducible from the seeded generator, so the manifests carry SHA-256 hashes instead. An archived snapshot carrying a permanent identifier will be deposited on acceptance.

file	contents
<i>Instance family</i>	
<code>gen_degeneracy_family.py</code>	generator and screening (section 2)
<code>degen_instances_paired/</code>	270 instances: <code>G_*.bin</code> , <code>Xstar_*.bin</code> , <code>manifest.tsv</code> (SHA-256), <code>cases.csv</code> , <code>seed_screening.csv</code>
<code>materialize_exact_ystars.py</code>	exact $y^* = -\text{diag}(W)$ per instance
<code>degen_instances_n500_paired/</code>	270 instances at $n = 500$, same grid, all used in section 5.6
<i>Solvers</i>	
<code>bench_sbb_dual.jl</code>	<code>sbb_dual</code> , <code>newton_ncm_globalized</code> , <code>agd_sdaj_ncm</code> , <code>dykstra_apm</code> , <code>snap!</code> trajectory recorder, <code>run_ranking_study</code>
<code>agd_sdaj.jl</code>	standalone AGD-SDAJ used for verification (section B)
<code>export_agd_verify_matrices.jl</code>	exports P7/P8/P9 from Anymatrix CORRINV
<i>Results and analysis</i>	
<code>ranking_kkt270_v2.csv</code>	one row per EVD, five solvers, 270 instances
<code>analyze_ranking.py</code> , <code>ranking_kkt270_v2_analysis.md</code>	tables 5 to 7, 9 and 12 to 15
<code>timing_kkt270_{native,matched}.csv</code>	single-threaded timings, six variants
<code>analyze_timing.py</code> , <code>timing_kkt270_analysis.md</code>	tables 2, 3 and 16
<code>globalization_*.csv</code> , <code>globalization_correction_report.md</code>	Newton Armijo stall analysis (section 4.2)

Figure data. The files `figures/data/*.dat` in this manuscript’s source were extracted from the CSVs above. The file `evd_vs_eps.dat` applies the cost-to-target rule of `analyze_ranking.py` (accepted-only, reach and hold) again, on a half-decade grid of ε . At the five decade targets it reproduces table 6, up to that table’s rounding of half-integer medians.

References

- [1] Neculai Andrei. An acceleration of gradient descent algorithm with backtracking for unconstrained optimization. *Numerical Algorithms*, 42:63–73, 2006.
- [2] Nicolas F. Armijo, Yunier Bello-Cruz, and Gabriel Haeser. A semi-smooth Newton method for general projection equations applied to the nearest correlation matrix problem. *Optimization*, 75:2293–2315, 2026. Published online 2025.
- [3] Vahid Beiranvand, Warren Hare, and Yves Lucet. Best practices for comparing optimization algorithms. *Optimization and Engineering*, 18(4):815–848, 2017.
- [4] Rüdiger Borsdorf. A Newton algorithm for the nearest correlation matrix. MSc thesis, University of Manchester, 2007. MIMS EPrint 1085.
- [5] Rüdiger Borsdorf and Nicholas J. Higham. A preconditioned Newton algorithm for the nearest correlation matrix. *IMA Journal of Numerical Analysis*, 30(1):94–107, 2010.
- [6] Rüdiger Borsdorf, Nicholas J. Higham, and Marcos Raydan. Computing a nearest correlation matrix with factor structure. *SIAM Journal on Matrix Analysis and Applications*, 31(5):2603–2622, 2010.
- [7] Stephen Boyd and Lin Xiao. Least-squares covariance matrix adjustment. *SIAM Journal on Matrix Analysis and Applications*, 27(2):532–546, 2005.
- [8] Paul H. Calamai, Luís N. Vicente, and Joaquim J. Júdice. A new technique for generating quadratic programming test problems. *Mathematical Programming*, 61:215–231, 1993.
- [9] Elizabeth D. Dolan and Jorge J. Moré. Benchmarking optimization software with performance profiles. *Mathematical Programming*, 91(2):201–213, 2002.
- [10] Elizabeth D. Dolan, Jorge J. Moré, and Todd S. Munson. Optimality measures for performance profiles. *SIAM Journal on Optimization*, 16(3):891–909, 2006.
- [11] Nicholas Gould and Jennifer Scott. A note on performance profiles for benchmarking software. *ACM Transactions on Mathematical Software*, 43(2):15:1–15:5, 2016.
- [12] William W. Hager and Hongchao Zhang. A new conjugate gradient method with guaranteed descent and an efficient line search. *SIAM Journal on Optimization*, 16(1):170–192, 2005.
- [13] Ernst Hairer and Gerhard Wanner. *Solving Ordinary Differential Equations II: Stiff and Differential-Algebraic Problems*. Springer Series in Computational Mathematics. Springer, Berlin, Heidelberg, 2 edition, 1996.
- [14] Nikolaus Hansen, Anne Auger, Raymond Ros, Olaf Mersmann, Tea Tušar, and Dimo Brockhoff. COCO: a platform for comparing continuous optimizers in a black-box setting. *Optimization Methods and Software*, 36(1):114–144, 2021.

- [15] Nicholas J. Higham. Computing the nearest correlation matrix—a problem from finance. *IMA Journal of Numerical Analysis*, 22(3):329–343, 2002.
- [16] Nicholas J. Higham and Mantas Mikaitis. Anymatrix: an extensible MATLAB matrix collection. *Numerical Algorithms*, 90:1175–1196, 2022.
- [17] Nicholas J. Higham and Nataša Strabić. Anderson acceleration of the alternating projections method for computing the nearest correlation matrix. *Numerical Algorithms*, 72:1021–1042, 2016.
- [18] Duc Quoc Huynh and Feng-Nan Hwang. Accelerated gradient descent with quasi-Newton preconditioning for nearest correlation matrix problems. *Numerical Linear Algebra with Applications*, 32(6):e70045, 2025.
- [19] Oyeyinka E. Ibrahim. A safeguarded spectral gradient method for the nearest correlation matrix. Manuscript in preparation, 2026.
- [20] Melanie L. Lenard and Michael Minkoff. Randomly generated test problems for positive definite quadratic programming. *ACM Transactions on Mathematical Software*, 10(1):86–96, 1984.
- [21] Dong-Hui Li and Masao Fukushima. A derivative-free line search and global convergence of Broyden-like method for nonlinear equations. *Optimization Methods and Software*, 13(3):181–201, 2000.
- [22] Jérôme Malick. A dual approach to semidefinite least-squares problems. *SIAM Journal on Matrix Analysis and Applications*, 26(1):272–284, 2004.
- [23] Jorge J. Moré, Burton S. Garbow, and Kenneth E. Hillstom. Testing unconstrained optimization software. *ACM Transactions on Mathematical Software*, 7(1):17–41, 1981.
- [24] Jorge J. Moré and Stefan M. Wild. Benchmarking derivative-free optimization algorithms. *SIAM Journal on Optimization*, 20(1):172–191, 2009.
- [25] Houduo Qi and Defeng Sun. A quadratically convergent Newton method for computing the nearest correlation matrix. *SIAM Journal on Matrix Analysis and Applications*, 28(2):360–385, 2006.
- [26] Roman Vershynin. *High-Dimensional Probability: An Introduction with Applications in Data Science*. Cambridge University Press, second edition, 2026.
- [27] Hua Wei and Henry Wolkowicz. Generating and measuring instances of hard semidefinite programs. *Mathematical Programming*, 125(1):31–45, 2010.

Algorithm 1 Newton-SIN-BH (semismooth Newton, cached factorization, BH line search)

Require: G , tolerance τ ; $\sigma = 10^{-4}$, $\mu = 0.9$, $\gamma = 100$, $u = 2^{-53}$

```

1:  $y \leftarrow 0$ ;  $(\theta, g, X, \Lambda, P) \leftarrow \text{EVAL}(y)$  ▷ keep the factorization
2: while  $\|g\| > \tau$  do
3:   Solve  $(V + 10^{-12}I)d = -g$  by CG, with  $V$  the generalized Jacobian from the cached  $(\Lambda, P)$ ;
   stop at  $\|r\| \leq \min\{0.1, \|g\|\}\|g\|$  or on negative curvature
4:   if  $d$  not finite or  $g^\top d \geq 0$  then  $d \leftarrow -g$ 
5:   end if
6:    $\alpha \leftarrow 1$ 
7:   loop
8:      $(\theta_t, g_t, X_t, \Lambda_t, P_t) \leftarrow \text{EVAL}(y + \alpha d)$ 
9:     if  $\theta_t \leq \theta + \sigma \alpha g^\top d$  then accept; break
10:    end if
11:    if  $|\theta_t - \theta| < \gamma u(1 + |\theta_t| + |\theta|)$  then ▷ BH near-equality branch
12:      if full step has  $\|g_t^{(\alpha=1)}\| \leq (1 - \mu)\|g\|$  then accept the full step; break
13:      end if
14:      accept the unit steepest-descent step  $y - g$  (one further EVD); break
15:    end if
16:     $\alpha \leftarrow \alpha/2$ 
17:  end loop
18:  update  $(y, \theta, g, X, \Lambda, P)$  from the accepted trial ▷ no recomputation
19: end while
20: return  $X$ 

```

Algorithm 2 AGD-SDAJ [18] (AGD-SDAJ-BH: with the marked branch)

Require: $G, \tau; q = 2, c = 10^{-4}, \rho = 0.5, \lambda_1 = \lambda_2 = 0.01, \epsilon_l = 10^{-5}, \epsilon_u = 10^8$

```
1:  $y \leftarrow 0; (\theta, g, X) \leftarrow \text{EVAL}(y)$ 
2: while  $\|g\| > \tau$  do
3:    $(y^p, \theta^p, g^p) \leftarrow (y, \theta, g); b \leftarrow e$  ▷  $B_0 = I$ , reset every outer iteration
4:   for  $i = 0, \dots, q - 1$  while  $\|g^p\| > \tau$  do ▷ QN-SDAJ inner steps
5:      $d \leftarrow -g^p \oslash b; \epsilon_i \leftarrow 1/(i+1)^2; \alpha \leftarrow 1$ 
6:     backtrack  $\alpha \leftarrow \rho\alpha$  until, with  $g_t = \nabla\theta(y^p + \alpha d)$ ,  

 $\|g_t\|^2 \leq (1 + \epsilon_i)\|g^p\|^2 - \lambda_1\alpha^2\|g^p\|^2 - \lambda_2\alpha^2\|d\|^2$ 
    ▷ eq. (6)
7:      $\beta \leftarrow a/b'$  with  $a = -g^{p\top}d, b' = (g_t - g^p)^\top d$  if  $b' > 0$ , else  $\beta \leftarrow 1$  ▷ general direction
8:      $s \leftarrow \beta\alpha d; y^+ \leftarrow y^p + s$ ; evaluate at  $y^+$  (reuse the trial if  $\beta = 1$ )
9:      $b_j \leftarrow z_j/s_j$  if  $z_j/s_j \in [\epsilon_l, \epsilon_u]$ , else 1, where  $z = g^+ - g^p$ 
10:     $(y^p, \theta^p, g^p) \leftarrow (y^+, \theta^+, g^+)$ 
11:  end for
12:  if  $\theta^p > \theta$  then  $(y^p, \theta^p, g^p) \leftarrow (y, \theta, g)$ 
13:  end if
14:  if  $\|g^p\| \leq \tau$  then  $y \leftarrow y^p$ ; break ▷ early exit at  $x^{\text{pre}}$ 
15:  end if
16:   $d \leftarrow -g^p; \alpha \leftarrow 1$ 
17:  loop ▷ Algorithm 1, line 5
18:     $(\theta_t, g_t) \leftarrow \text{EVAL}(y^p + \alpha d)$ 
19:    if  $\theta_t \leq \theta^p + c\alpha g^{p\top}d$  then break
20:    end if
21:    if BH and  $|\theta_t - \theta^p| < \gamma u(1 + |\theta_t| + |\theta^p|)$  and  $\|g_t\| < \|g^p\|$  then break
22:    end if
23:     $\alpha \leftarrow \rho\alpha$ 
24:  end loop
25:   $\beta \leftarrow \|g^p\|^2 / ((g^p - g_t)^\top g^p)$  if the denominator is  $> 0$ , else 1
26:   $y \leftarrow y^p - \beta\alpha g^p; (\theta, g, X) \leftarrow \text{EVAL}(y)$  (reuse the trial if  $\beta = 1$ )
27: end while
28: return  $X$ 
```

Algorithm 3 SBB-Dual (clamped spectral gradient on the dual)

Require: $G, \tau, \eta = 0.01$

```
1:  $y \leftarrow 0; (\theta, g, X) \leftarrow \text{EVAL}(y)$ 
2: while  $\|g\| > \tau$  do
3:   if first update, or  $s^\top w \leq 0$ , where  $s = y - y_{\text{prev}}, w = g - g_{\text{prev}}$  then
4:      $t \leftarrow 1$ 
5:   else
6:      $t \leftarrow \min\{\max\{s^\top s / s^\top w, \eta\}, 2 - \eta\}$  ▷ clamped BB1
7:   end if
8:    $(y_{\text{prev}}, g_{\text{prev}}) \leftarrow (y, g); y \leftarrow y - tg; (\theta, g, X) \leftarrow \text{EVAL}(y)$ 
9: end while
10:  $X \leftarrow D^{-1/2} X D^{-1/2}, D = \text{Diag}(\text{diag } X)$ ; set  $\text{diag}(X) = e$  ▷ rescaling epilogue
11: return  $X$ 
```

Algorithm 4 Dykstra-APM (alternating projections with Dykstra correction)

Require: G, τ

1: $X \leftarrow G; \Delta S \leftarrow 0; \Delta U \leftarrow 0$

2: **loop**

3: $Y \leftarrow \Pi_{\mathcal{S}_+^n}(X + \Delta S)$

▷ one EVD

4: **if** $\|\text{diag}(Y) - e\|_2 \leq \tau$ **then return** Y

▷ PSD half-iterate

5: **end if**

6: $\Delta S \leftarrow X + \Delta S - Y$

7: $Z \leftarrow Y + \Delta U$ with $\text{diag}(Z) \leftarrow e; \Delta U \leftarrow Y + \Delta U - Z; X \leftarrow Z$

8: **end loop**
